


# Ethereum

## Introduction

Leander Jehl

# Ethereum

## Timeline



- 2014 Initial Ether sold on bitcoin
- 2016 DAO Attack: A Hacker exploited a smart contract bug, the community decided to undo attack, by discarding some blocks.
- 2020 Possible to stake ether and participate in PoS beacon chain, that votes on blocks created by PoW
- 2022 PoW deprecated and Ethereum using PoS to create blocks

# From PoW to PoS

# Ethereum - PoW

- Fast blocks (every 12sec)
- Structured P2P network
- Uses PoW hashing function not suited for ASICs
- Uncles: Blocks lost in a fork still get some reward


# Ethereum - PoS

- One staker (validator) needs to deposit 32 ether (>75.000\$ 2024)
- Stake can be slashed, but not increased  
-> 1.000.000 validators
- Validators need to run a node
- Can run one node for multiple validator ids.

**Problem:** Collecting signatures from  
1mill nodes takes a lot of time.

# Ethereum - PoS

## Slots, Epoch, and Time




- One epoch is 6.4 minutes
- Each validator needs to vote once in each epoch (vote -> reward, no vote -> get punished)
- Epoch divided in 32 slots (each 12 sec)
- Validators divided in 32 committees
- Vote in your slot!

One block every 12 sec without requiring one vote every 12 sec!

# Ethereum - PoS

## Proposers

- At the start of the epoch  
32 random validators (proposers)  
are selected to propose blocks
- Proposers also collect votes
- Proposers get extra reward



# Ethereum - Forks

Two mechanisms:

- **LMD Ghost:** Ensure new blocks can be created.
- **Casper FFG:** Create checkpoints
- Creating checkpoints is slow
- Forks can still arise between checkpoints




# Ethereum - Forks

## LMD Ghost

LMD Ghost is a fork choice rule (like longest chain):

- Rule which fork to extend




# Ethereum - Forks

## LMD Ghost

LMD Ghost: Extend fork with most votes

- Every validator can vote for one block

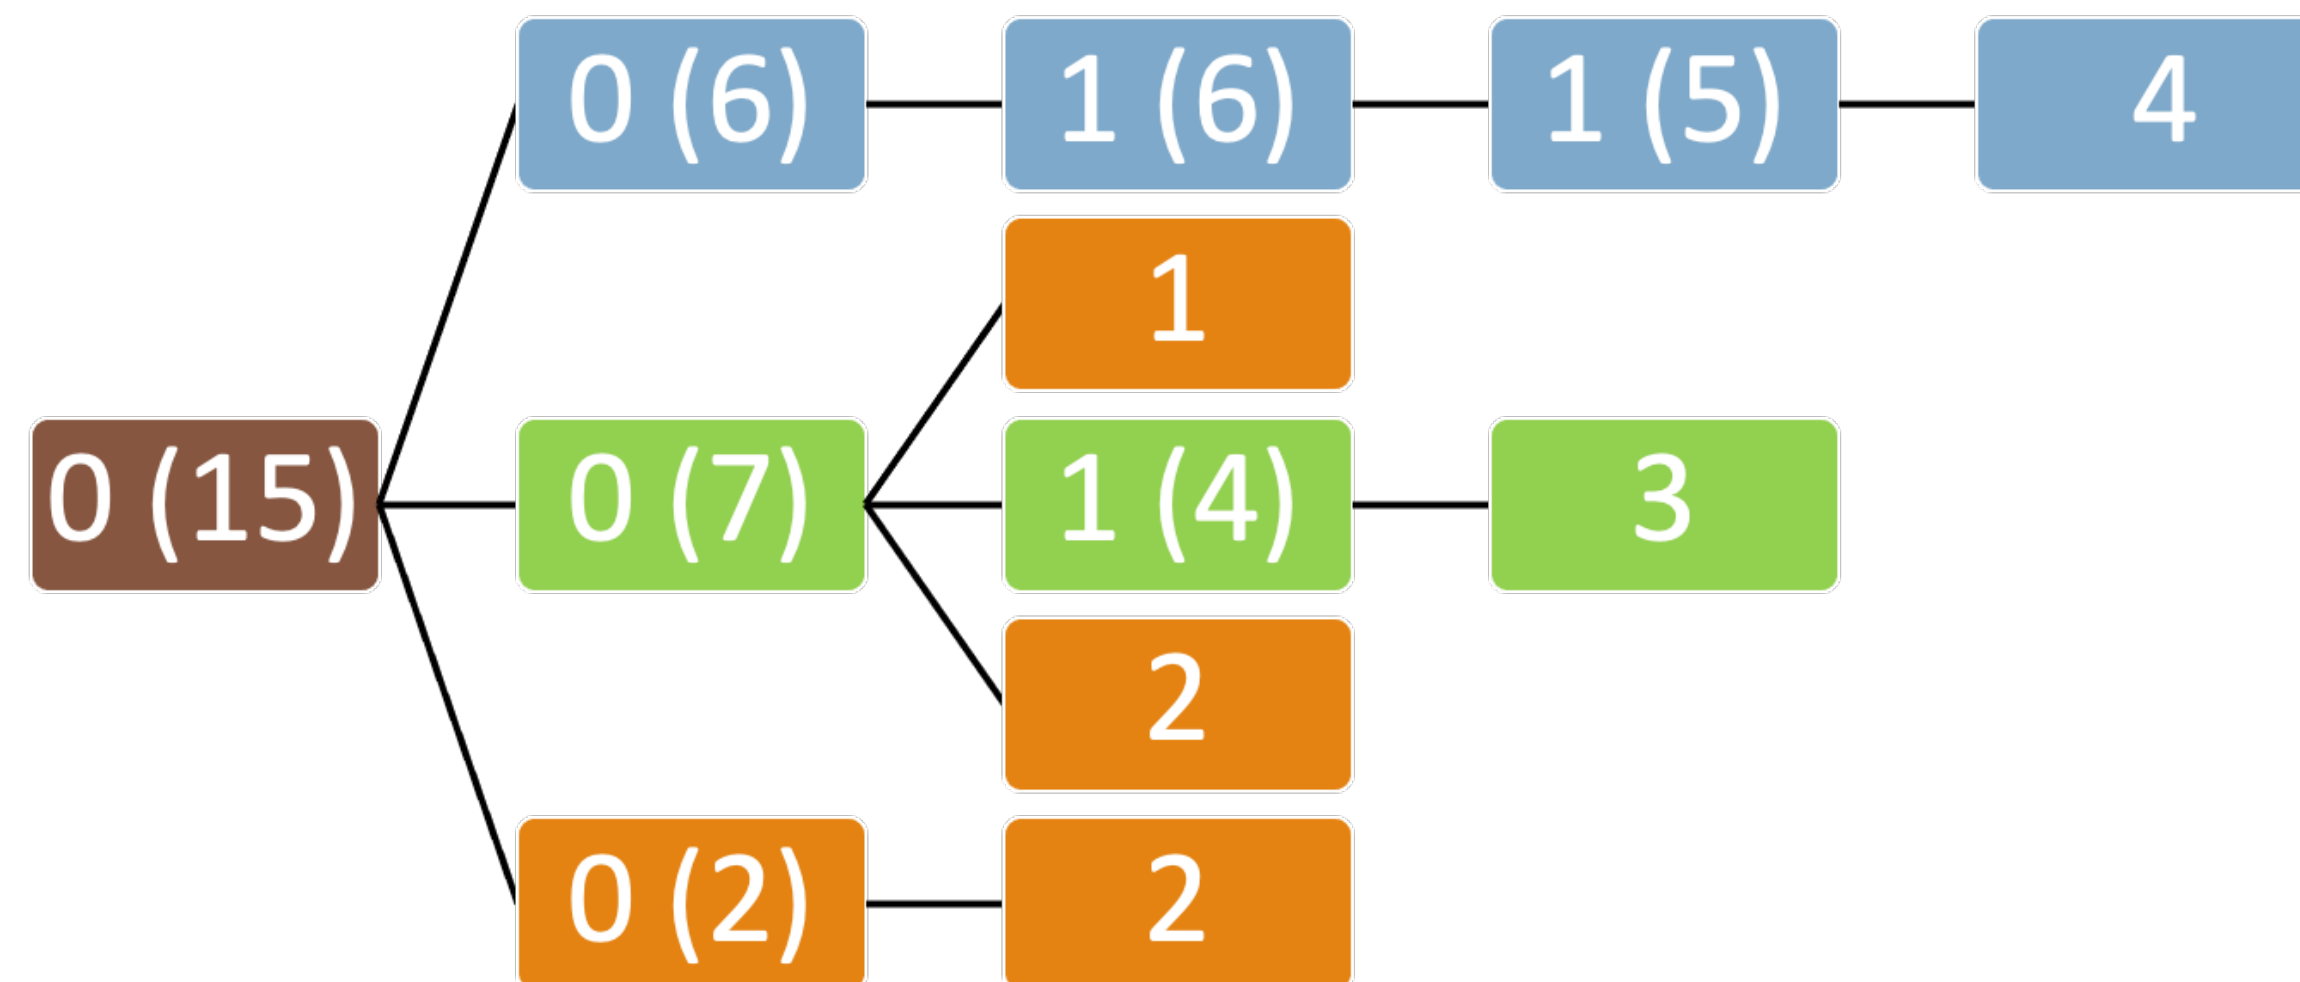


# Ethereum - Forks

## LMD Ghost

LMD Ghost: Extend fork with most votes

- Every validator can vote for one block
- Votes are summed for parents

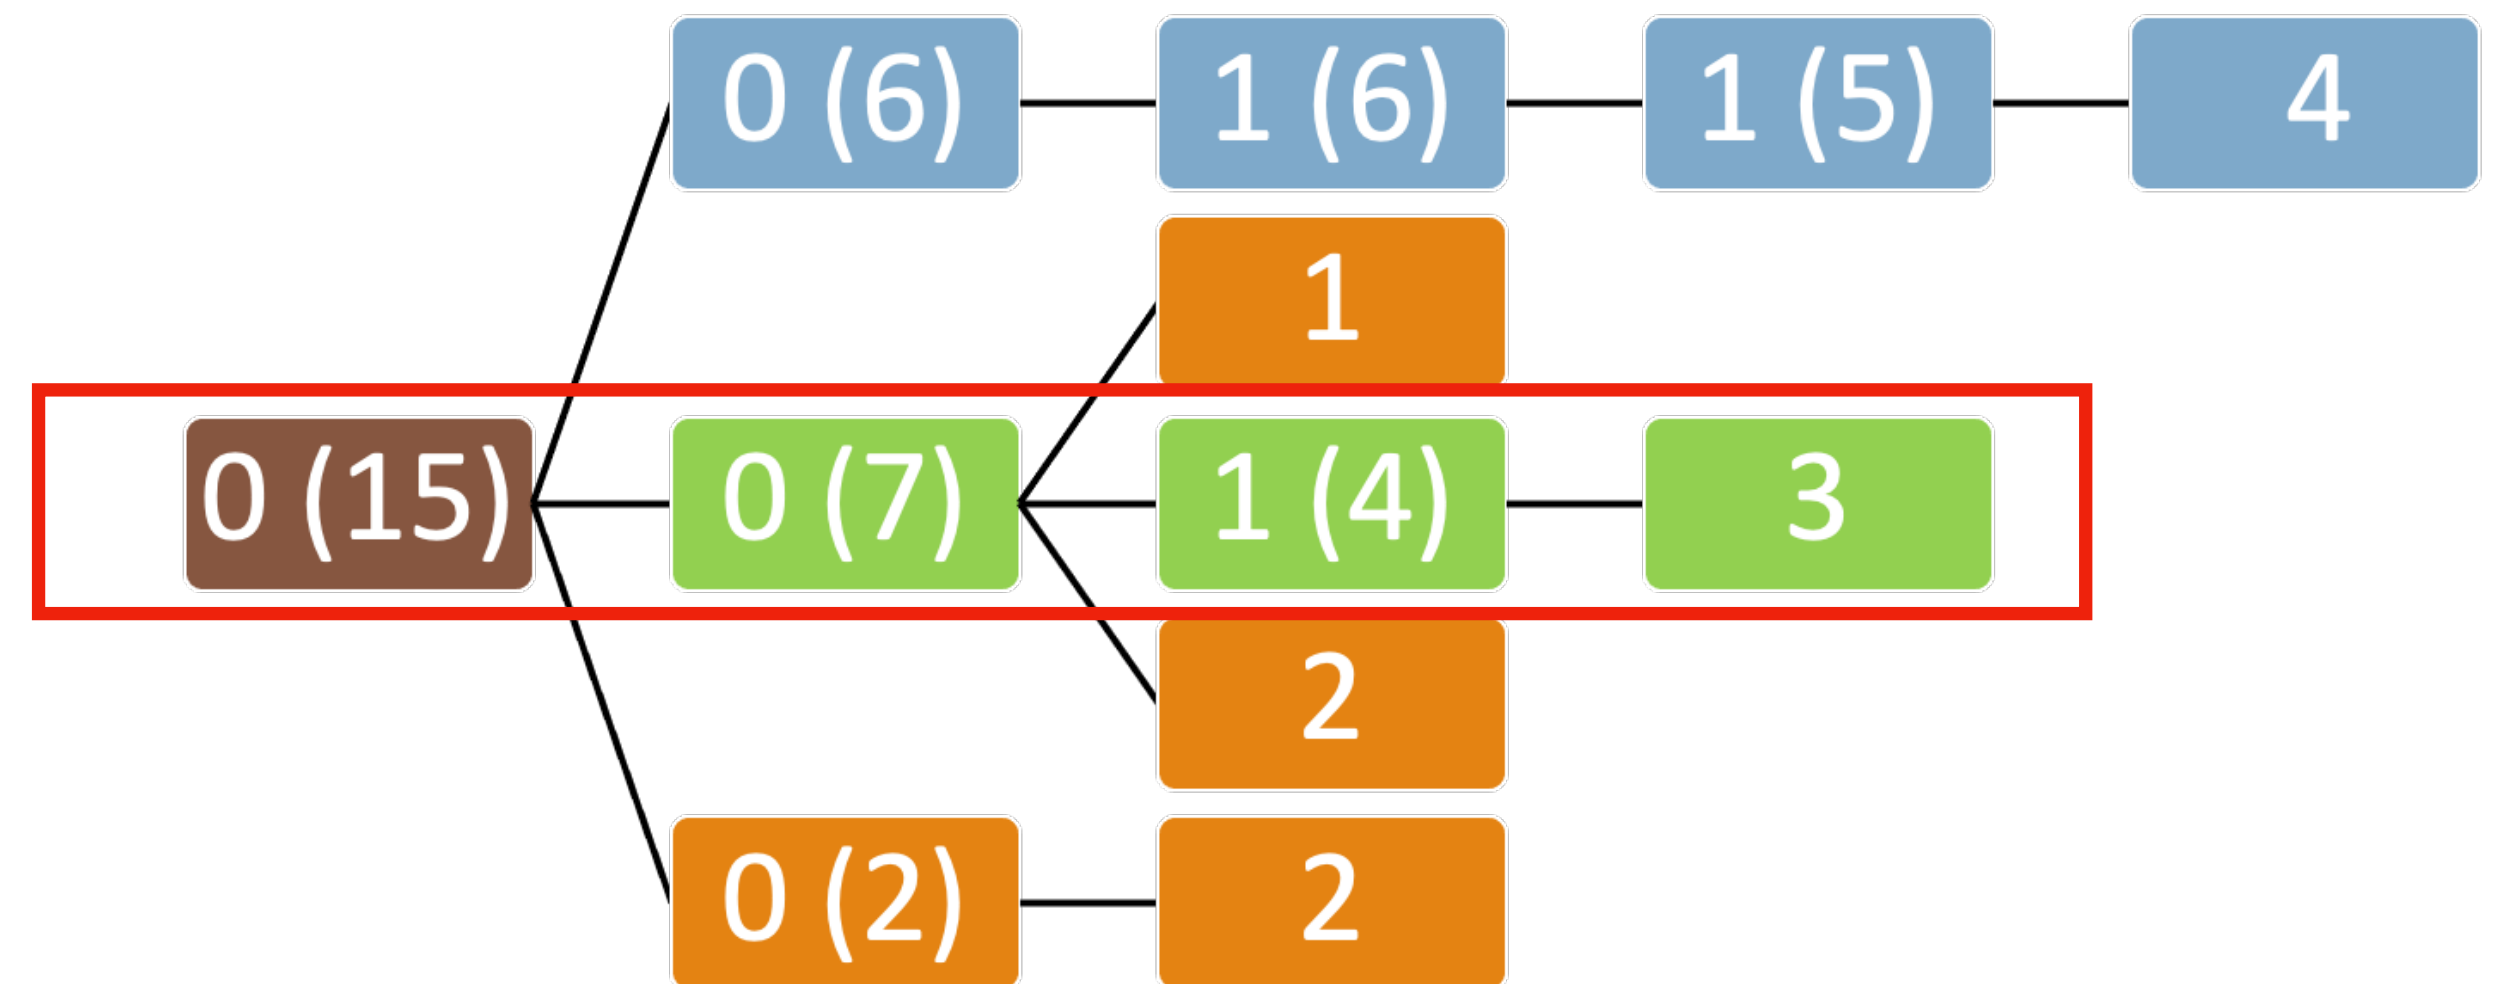


# Ethereum - Forks

## LMD Ghost

LMD Ghost: Extend fork with most votes

- Every validator can vote for one block
- Votes are summed for parents
- Choose fork with most votes

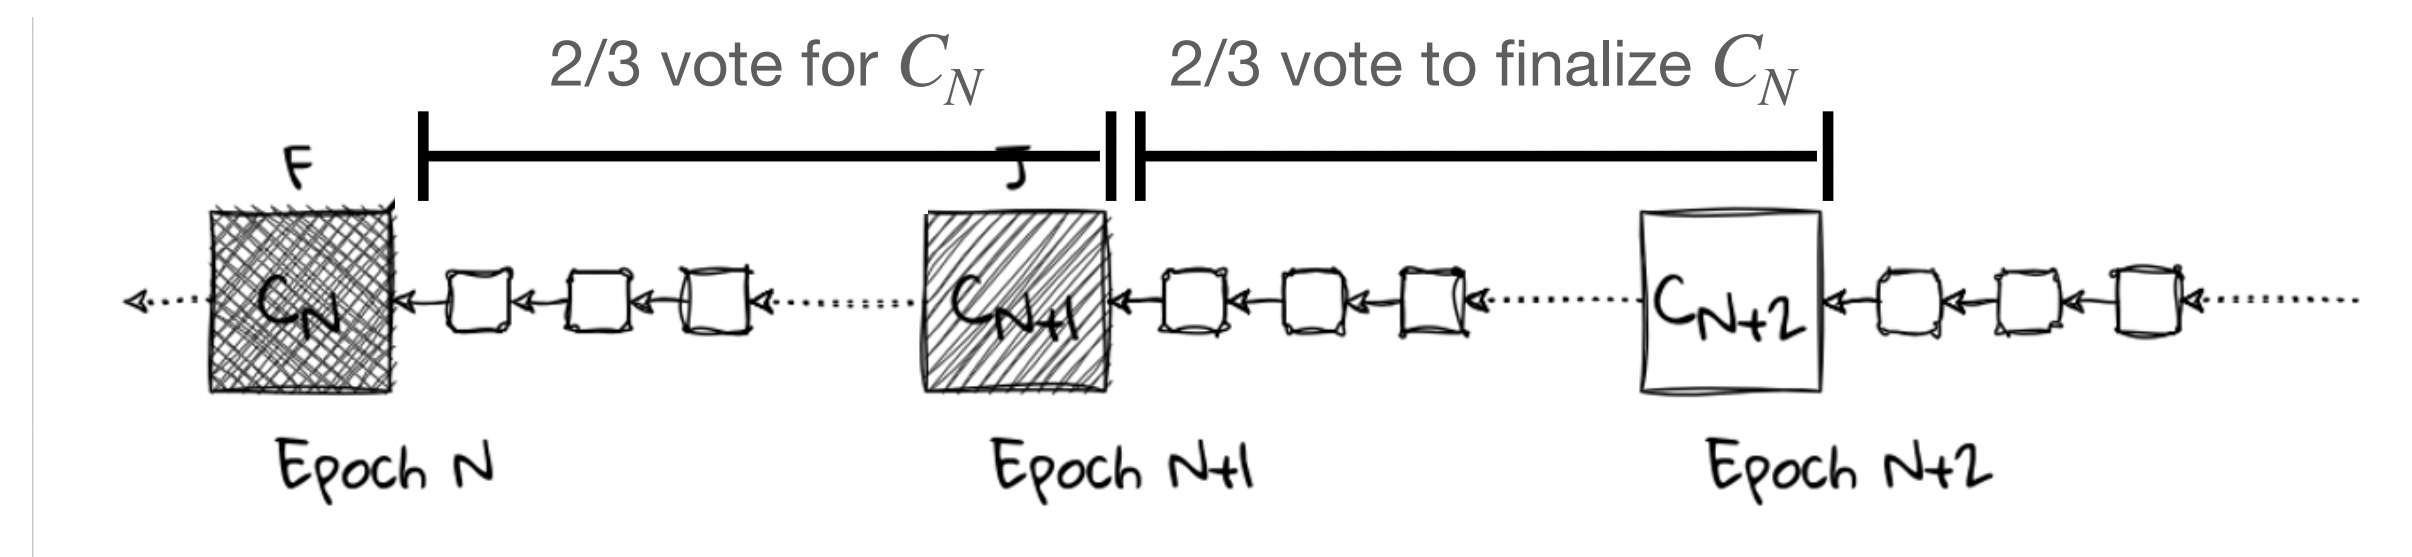


# Ethereum - Forks

## Casper FFG

Decide on checkpoint (consensus)

- Validators can vote to create checkpoint
- Need 2/3 of validators to vote for one checkpoint
- Needs to happen in 2 consecutive epochs




# Ethereum - Forks

## Casper FFG

Decide on checkpoint (consensus)

- Validators can vote to create checkpoint
- Need 2/3 of validators to vote for one checkpoint
- Needs to happen in 2 consecutive epochs



In fault free cases, creating a  
checkpoint takes 2 epochs  
12 minutes

# What are Smart Contracts

# Bitcoin scripts scripts

## Spending conditions

Transactions:

$$tx = \langle \underbrace{[(id_1, \sigma_1), (id_2, \sigma_2)]}_{\text{Inputs}}, \underbrace{[(pk_a, value_a), (pk_b, value_b)]}_{\text{Outputs}} \rangle$$



# Bitcoin scripts scripts

## Spending conditions

Transactions:

$$tx = \langle \underbrace{[(id_1, rd_1), (id_2, rd_2)]}_{\text{Inputs}}, \underbrace{[(s_a, value_a), (s_b, value_b)]}_{\text{Outputs}} \rangle$$

- $s_a$  a **spending condition**: output can be used if a value is supplied, that evaluates  $s_a$  to true
- $rd_1$  a **redeeming argument**: should ensure the script  $s_{id_1}$  returns true

# UTXO scripts

## Examples


Name	Spending condition	Redeeming argument	
P2Pk Pay to public key	Public key	Signature	
P2PkH Pay to public key hash	Hash of Public key	Public key and signature	
Multisig	$m$ public keys and parameter $k$	$k$ signatures	

# Ethereum

## Smart Contracts

In ethereum, a contract is like a object from OOP, with fields and methods

- variables containing state (stored in account, mutable)
- functions



# Ethereum

## Example: Simple Storage

compiler version

```
pragma solidity ^0.5.11;
```

contract

```
contract SimpleStorage {
```

```
    uint256 public storedData;
```

state

```
    function get() public view returns (uint256){  
        return storedData;  
    }
```

```
    function set(uint x_) public {  
        storedData = x_;  
    }
```

functions

```
}
```

# Ethereum

## Example: Simple Storage

Simple online IDE: <https://remix.ethereum.org/>

Fun tutorial: <https://cryptozombies.io/>

- Constructors
- Basic types and collections
- Visibility (private, public)
- Inheritance
- Modifiers (view, pure)

# Ethereum

## Example: Simple Storage

Who can invoke functions?

Who can view values?

Who can change the code?

```
pragma solidity ^0.5.11;

contract SimpleStorage {
    uint256 public storedData;

    function get() public view returns (uint256){
        return storedData;
    }

    function set(uint x_) public {
        storedData = x_;
    }
}
```

# Ethereum

## Example: Simple Storage

Who can invoke functions?

- any user

Who can view values?

- anyone

Who can change the code?

- noone

```
pragma solidity ^0.5.11;

contract SimpleStorage {
    uint256 public storedData;

    function get() public view returns (uint256){
        return storedData;
    }

    function set(uint x_) public {
        storedData = x_;
    }
}
```

# Ethereum

## Smart Contract code

Smart contract code is immutable and public

- Anyone can trust smart contract (if it is not too complex)
  - No need to trust the creator of the contract
- No one can fix bugs in the contract
- Anyone can find and exploit bugs in the contract



# Ethereum

## Smart Contract code

- Assembly for Ethereum Virtual machine (EVM)
- Compiled from higher level language (Solidity)
- Stored in account (codeHash)

# Ethereum

## Smart Contract

Why use a smart contract

-

# Ethereum

## Smart Contract

Why use a smart contract

- No legal system to enforce paper contract
- Cheaper than paper contract

# How does Ethereum enable Smart Contracts

# Ethereum

## Accounts

Ethereum uses accounts instead of UTXO.  
Thus the state of Ethereum contains for every account:

- **address:** e.g. pub-key hash
- **balance:** amount of Ether the address owns
- **nonce:** sequence number of last transaction sent from this account
- **storage root:** *only for non-user accounts (contract account)*
- **code hash:** *only for non-user accounts (contract account)*

# Ethereum

## Accounts

Smart Contracts are also represented as accounts.  
A contract account has:

- **address:** *e.g. hash from creator address & creation transaction nonce*
- **balance:** amount of Ether the address owns
- **nonce:** number of other contract created by this contract
- **storageRoot:** hash of data stored in this contract
- **codeHash:** hash of the code of this contract

Can create different contracts with  
the same code.  
(Like objects with the same type)

# Ethereum

## Accounts

In a contract written in Solidity, you can access:

- The address of the current contract:

```
address contractaddress = address(this);
```

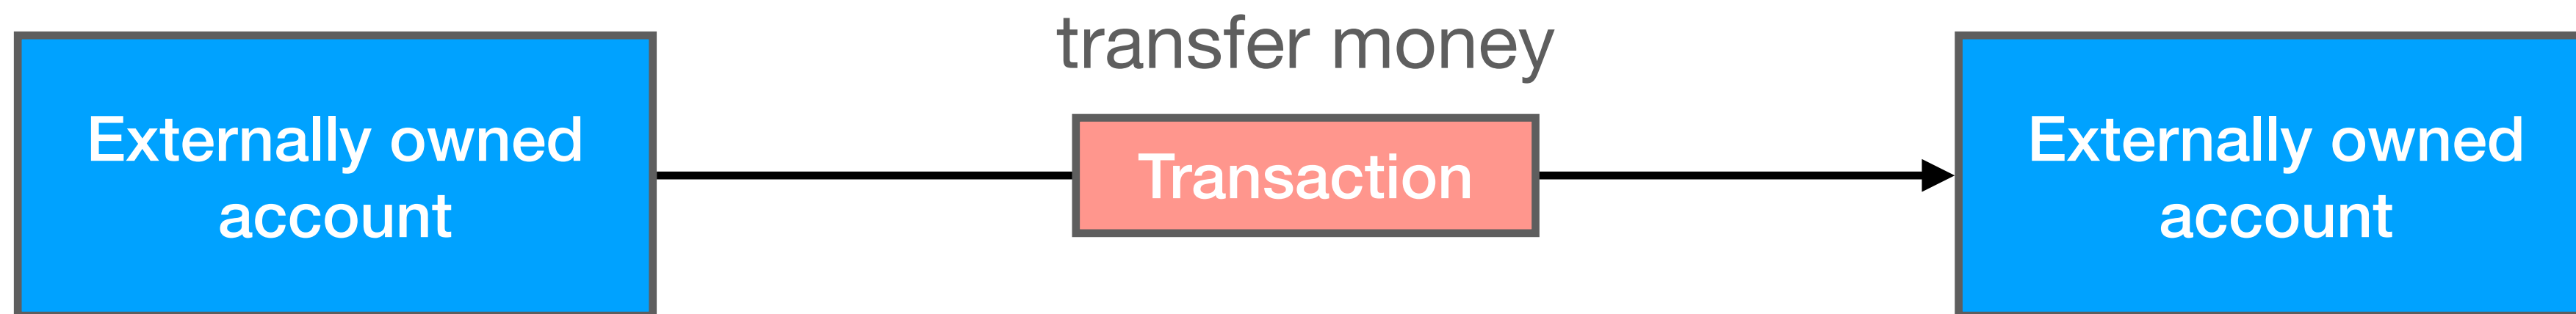
- The balance of the contract:

```
uint balance = contractaddress.balance;
```

# Ethereum

## Transactions and authentication

Transactions are used to transfer ether, invoke functions, and deploy new contracts.

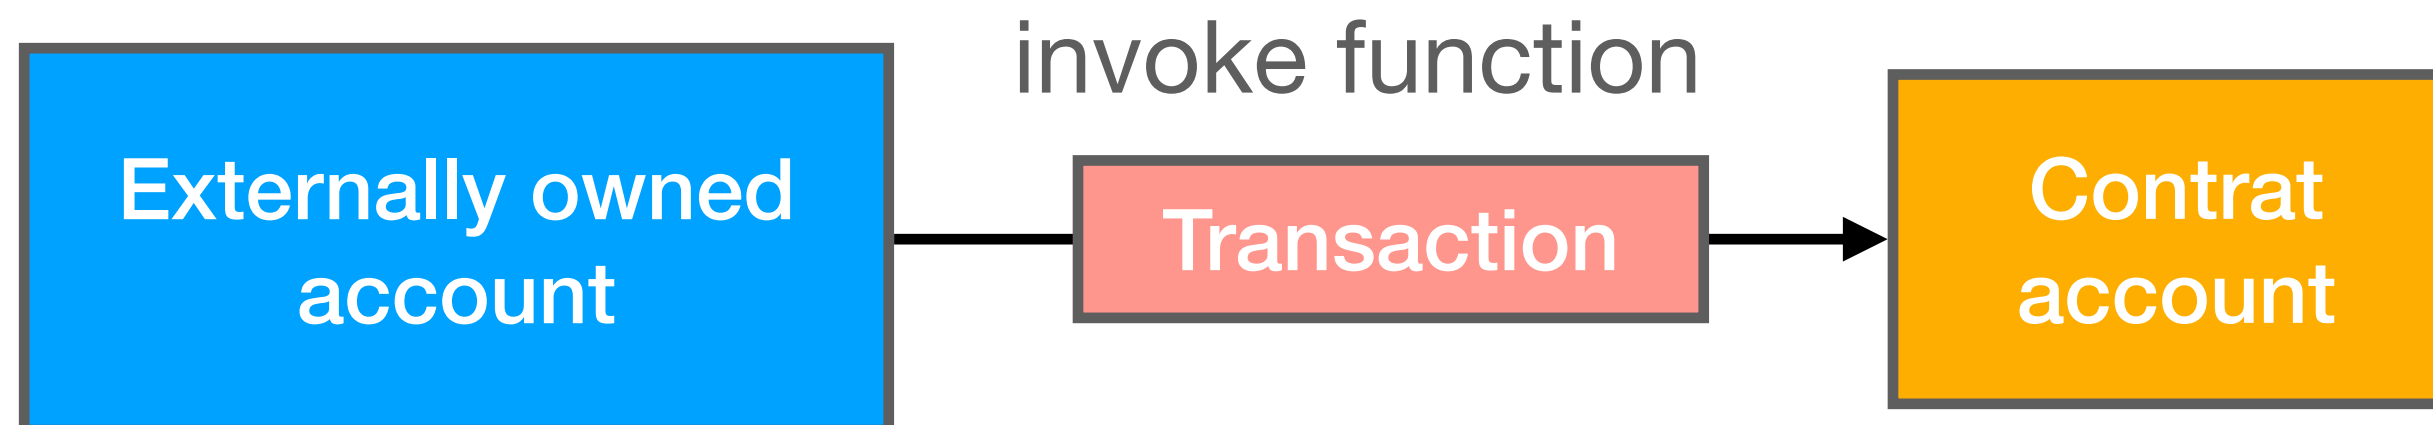




# Ethereum

## Transactions and authentication


Transactions are used to transfer ether, invoke functions, and deploy new contracts.



# Ethereum

## Transactions and authentication


Transactions are used to transfer ether, invoke functions, and deploy new contracts.



# Ethereum

## Transactions and authentication

Transactions are used to transfer ether, invoke functions, and deploy new contracts.



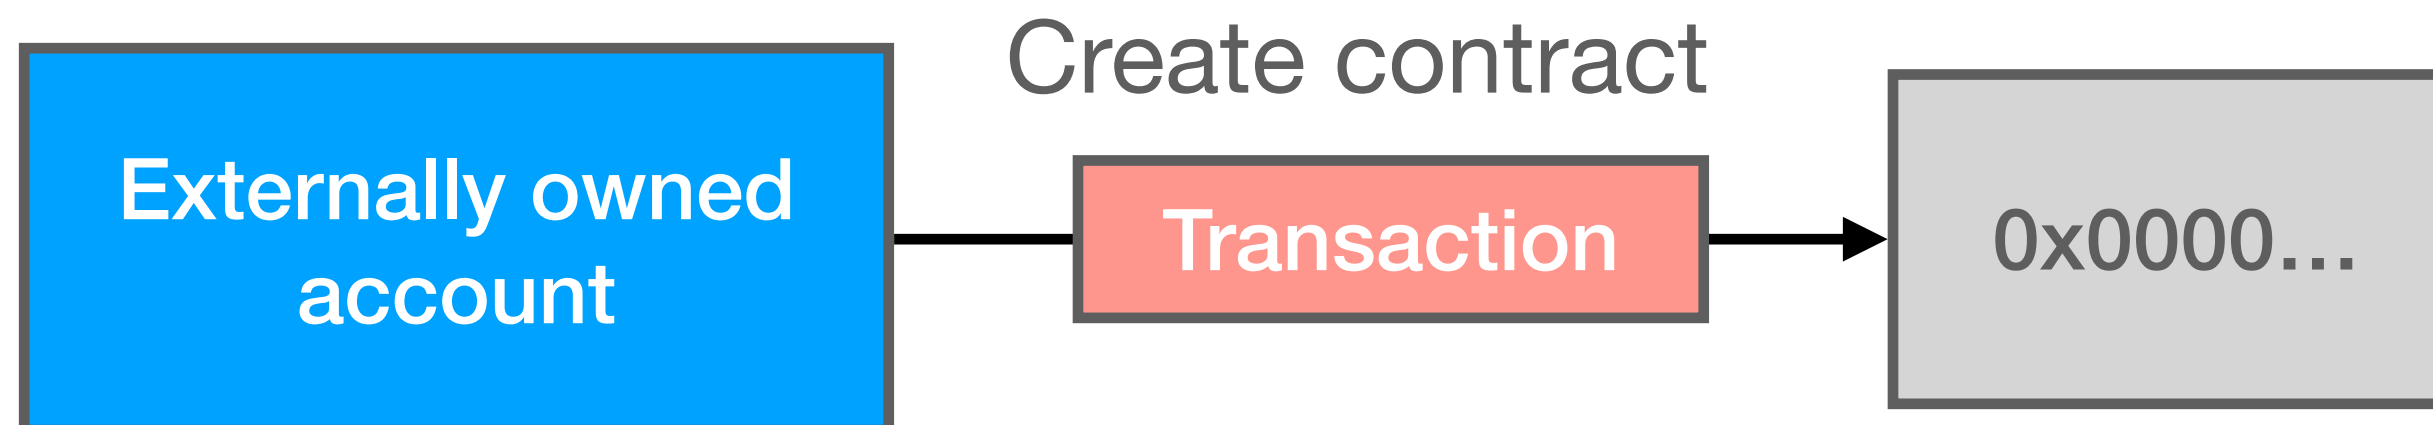
How does C1 know C2?

- C2 address supplied when deploying C1.
- C2 address supplied by transaction
- C2 address stored in contract state (can be updated)

# Ethereum

## Transactions and authentication


Transactions are used to transfer ether, invoke functions, and deploy new contracts.



# Ethereum

## Transactions and authentication

Transactions are used to transfer ether, invoke functions, and deploy new contracts.



# Ethereum

## Transactions and authentication

Transactions contain:

- *Nonce*: next sequence number for sender account
- *Gas price*: later
- *max gas*: later
- *Recipient*: destination Ethereum address
- *Value*: Amount of ether send to destination
- *Data*: Payload binary, e.g. function identifier and arguments
- *Signature*: Signature from sender, including his public key

# Ethereum

## Transaction validation

Transaction validation includes the following checks

- *Nonce*: is next sequence number for sender account
- Sender has sufficient balance to pay value and fees
- Transaction is correctly signed

When authenticating users in smart contract, we can rely on transaction validation!  
Use *msg.sender* to access address invoking transaction.

# Ethereum

## Solidity example

```
contract SimpleBank {
    mapping(address => uint) private balances;
    address public owner;

    // function SimpleBank() deprecated syntax for
    constructor() public {
        owner = msg.sender;
    }

    function deposit() public payable returns(uint) {
        balances[msg.sender] += msg.value;
        return balances[msg.sender];
    }

    function withdraw(uint withdrawAmount) public returns (uint remainingBal){
        if (balances[msg.sender] >= withdrawAmount){
            balances[msg.sender] -= withdrawAmount;
            // this throws an error if fails.
            msg.sender.transfer(withdrawAmount);
        }
        return balances[msg.sender];
    }

    function balance() view public returns (uint) {
        return balances[msg.sender];
    }
}
```



# Ethereum

## Solidity example

What happens if data is empty?

- Money transfered to account. Default function run.

What is *msg.sender* for internal transactions?

- address of sending contract
  - a contract can have money in our bank!

# Ethereum

## Solidity exceptions

If a smart contract throws an exception, or error, state is reverted.

# Ethereum

## Solidity view functions

View functions are read only

- Do not require fees
- Read contract state at one node

How can you know that a read  
from one node is correct?

**State stored in Ethereum  
blockchain**

# Bitcoin

## Block structure

### Header:

PrevBlockhash  
Nonce  
Timestamp

### Transaction data

Merkle tree

Merkle tree allows to easily proof that a transaction is included in a block.

State of the blockchain (UTXO)  
is not in the block.

# Ethereum

## Block structure


### Header:

PrevBlockhash  
Nonce  
Timestamp

State root hash  
Receipts root hash

### Transaction data

Merkle tree



# Ethereum

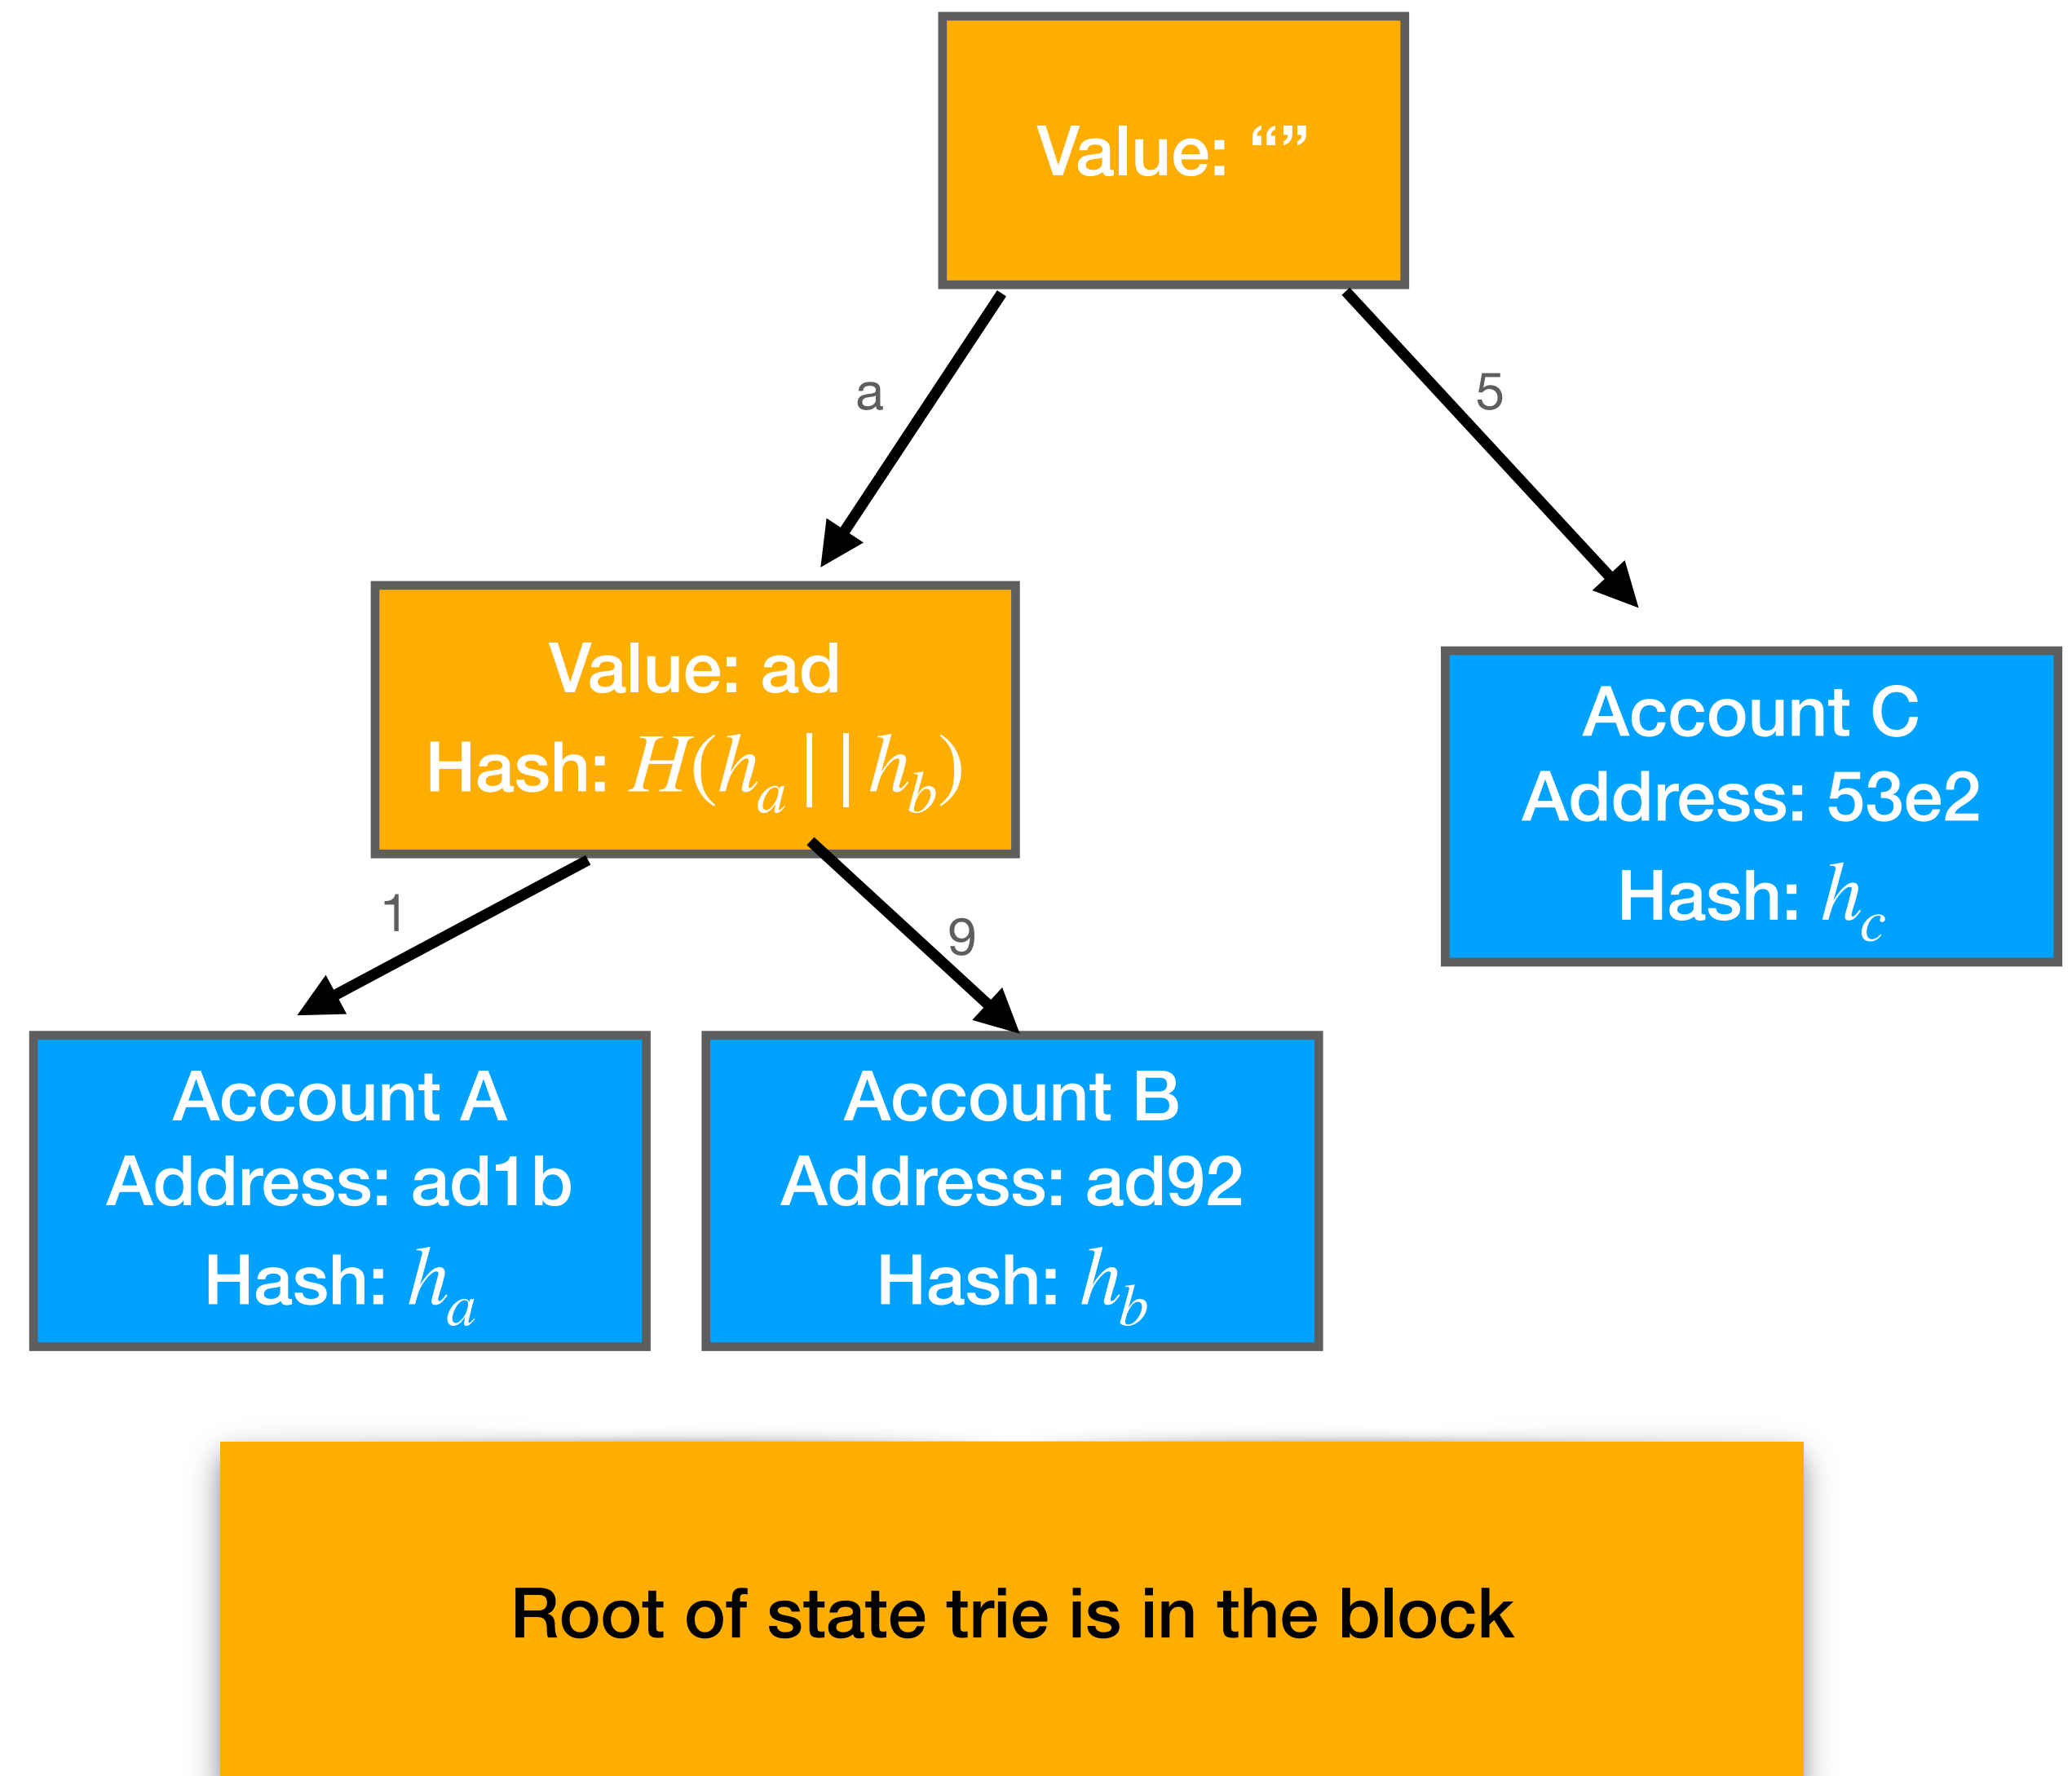
## State trie

Stores accounts:

Address:  
[Value,  
Nonce,  
StorageRoot,  
CodeHash]

Trie:  
Merkle tree that supports

update  
lookup  
proof



# Ethereum

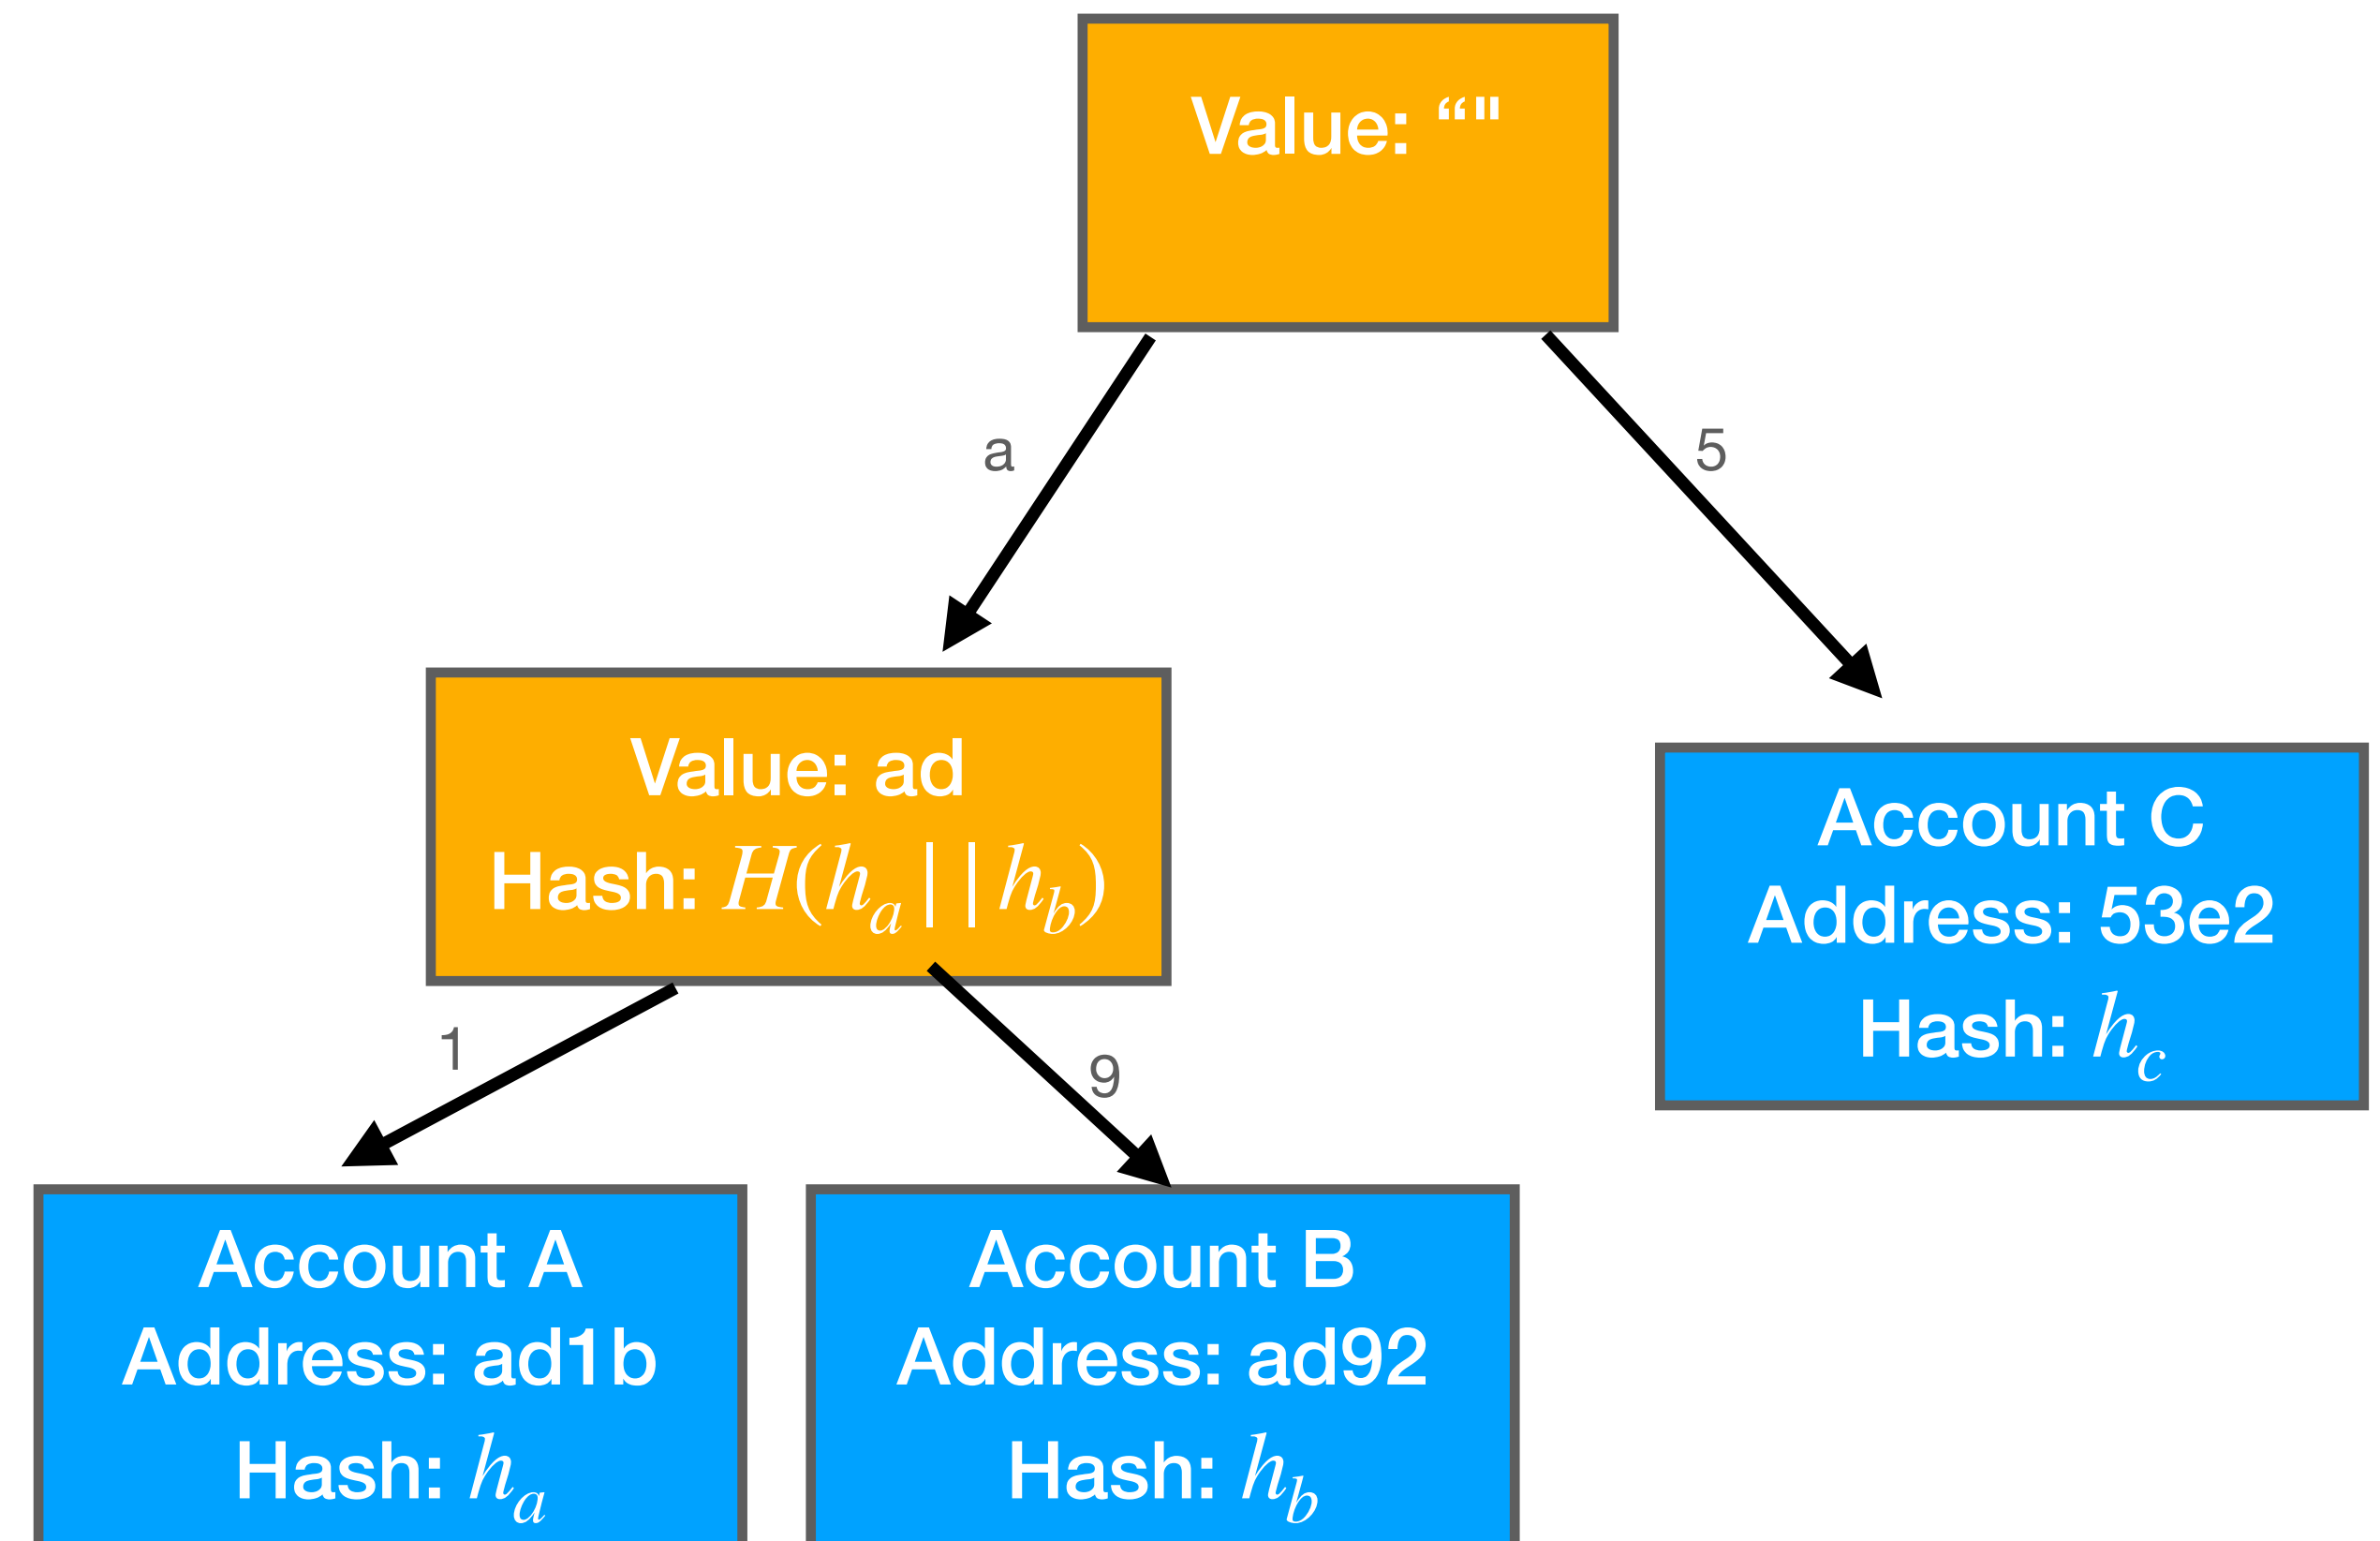
## State trie

Stores accounts:

Address:  
[Value,  
Nonce,  
StorageRoot,  
CodeHash]

Trie:  
Merkle tree that supports

update  
lookup  
proof



SC Variable changes

- > SC storage root changes
- > SC Account hash changes
- > State trie root changes



# Ethereum


## State trie

Stores accounts:

Address:  
[Value,  
Nonce,  
StorageRoot,  
CodeHash]

Trie:  
Merkle tree that supports

update  
lookup  
proof



StorageRoot is the root of a different trie.

# Ethereum

## State trie

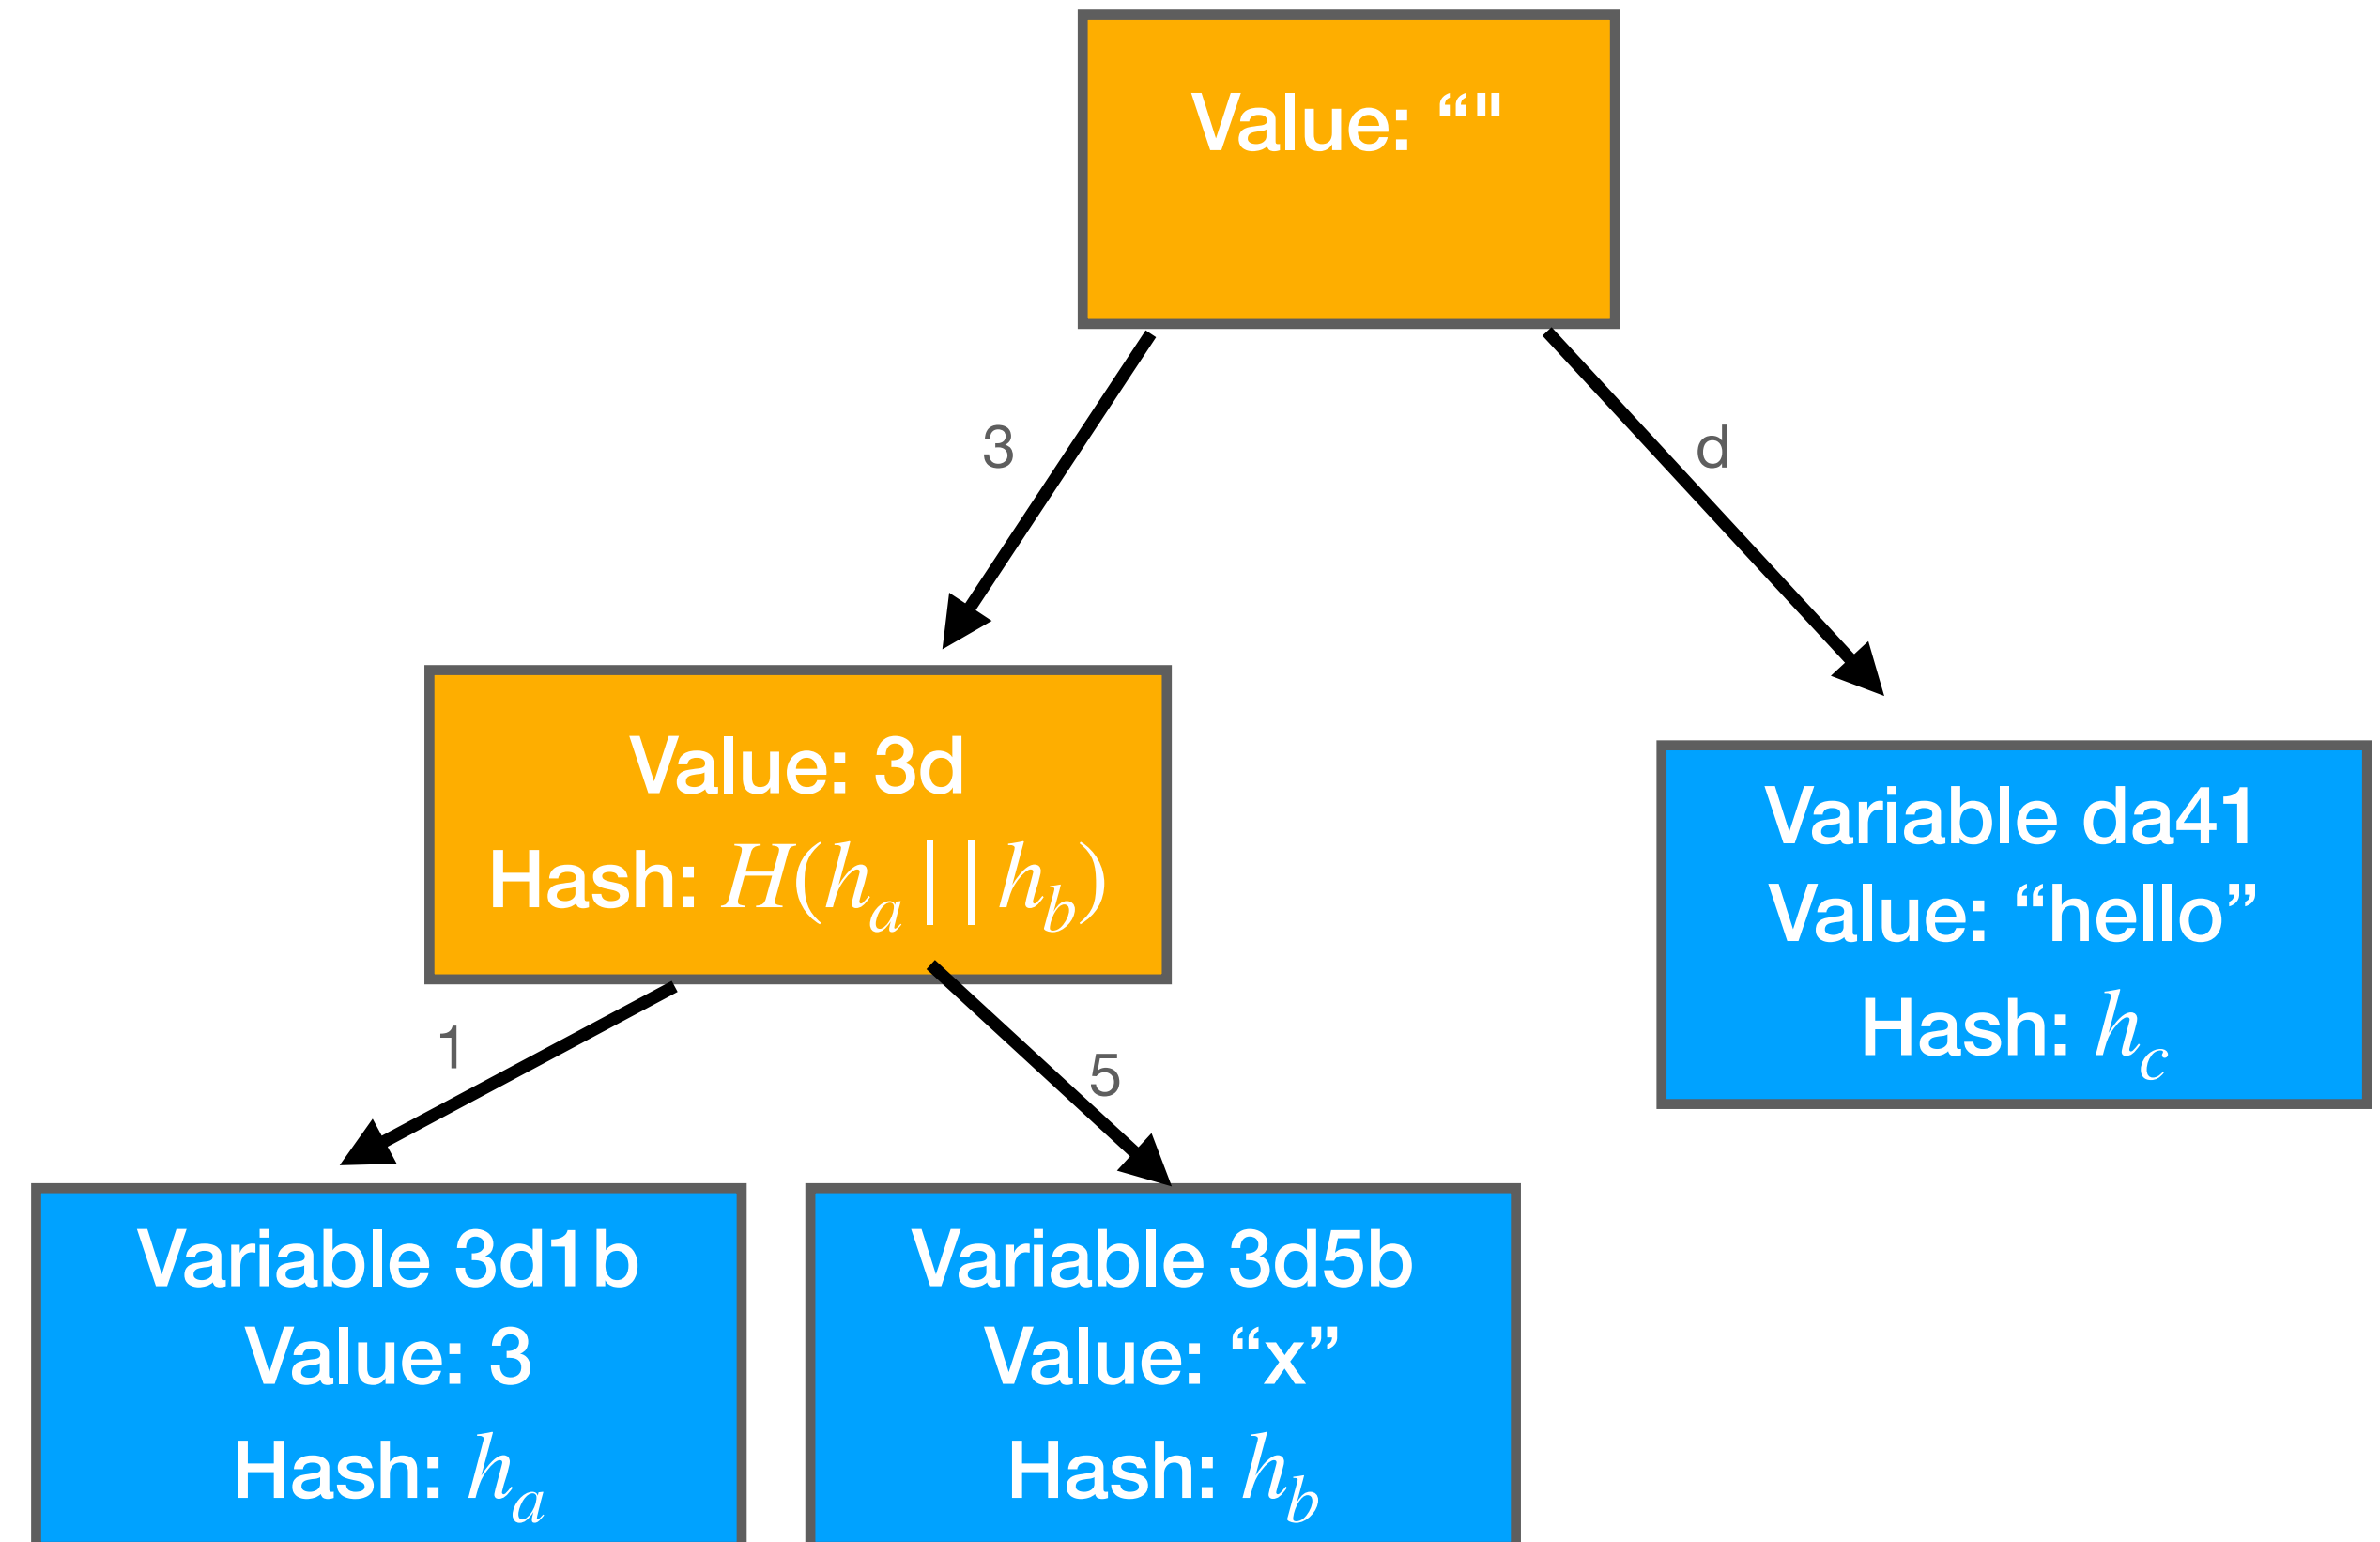
Stores accounts:

Address:  
[Value,  
Nonce,  
StorageRoot,  
CodeHash]

Trie:  
Merkle tree that supports

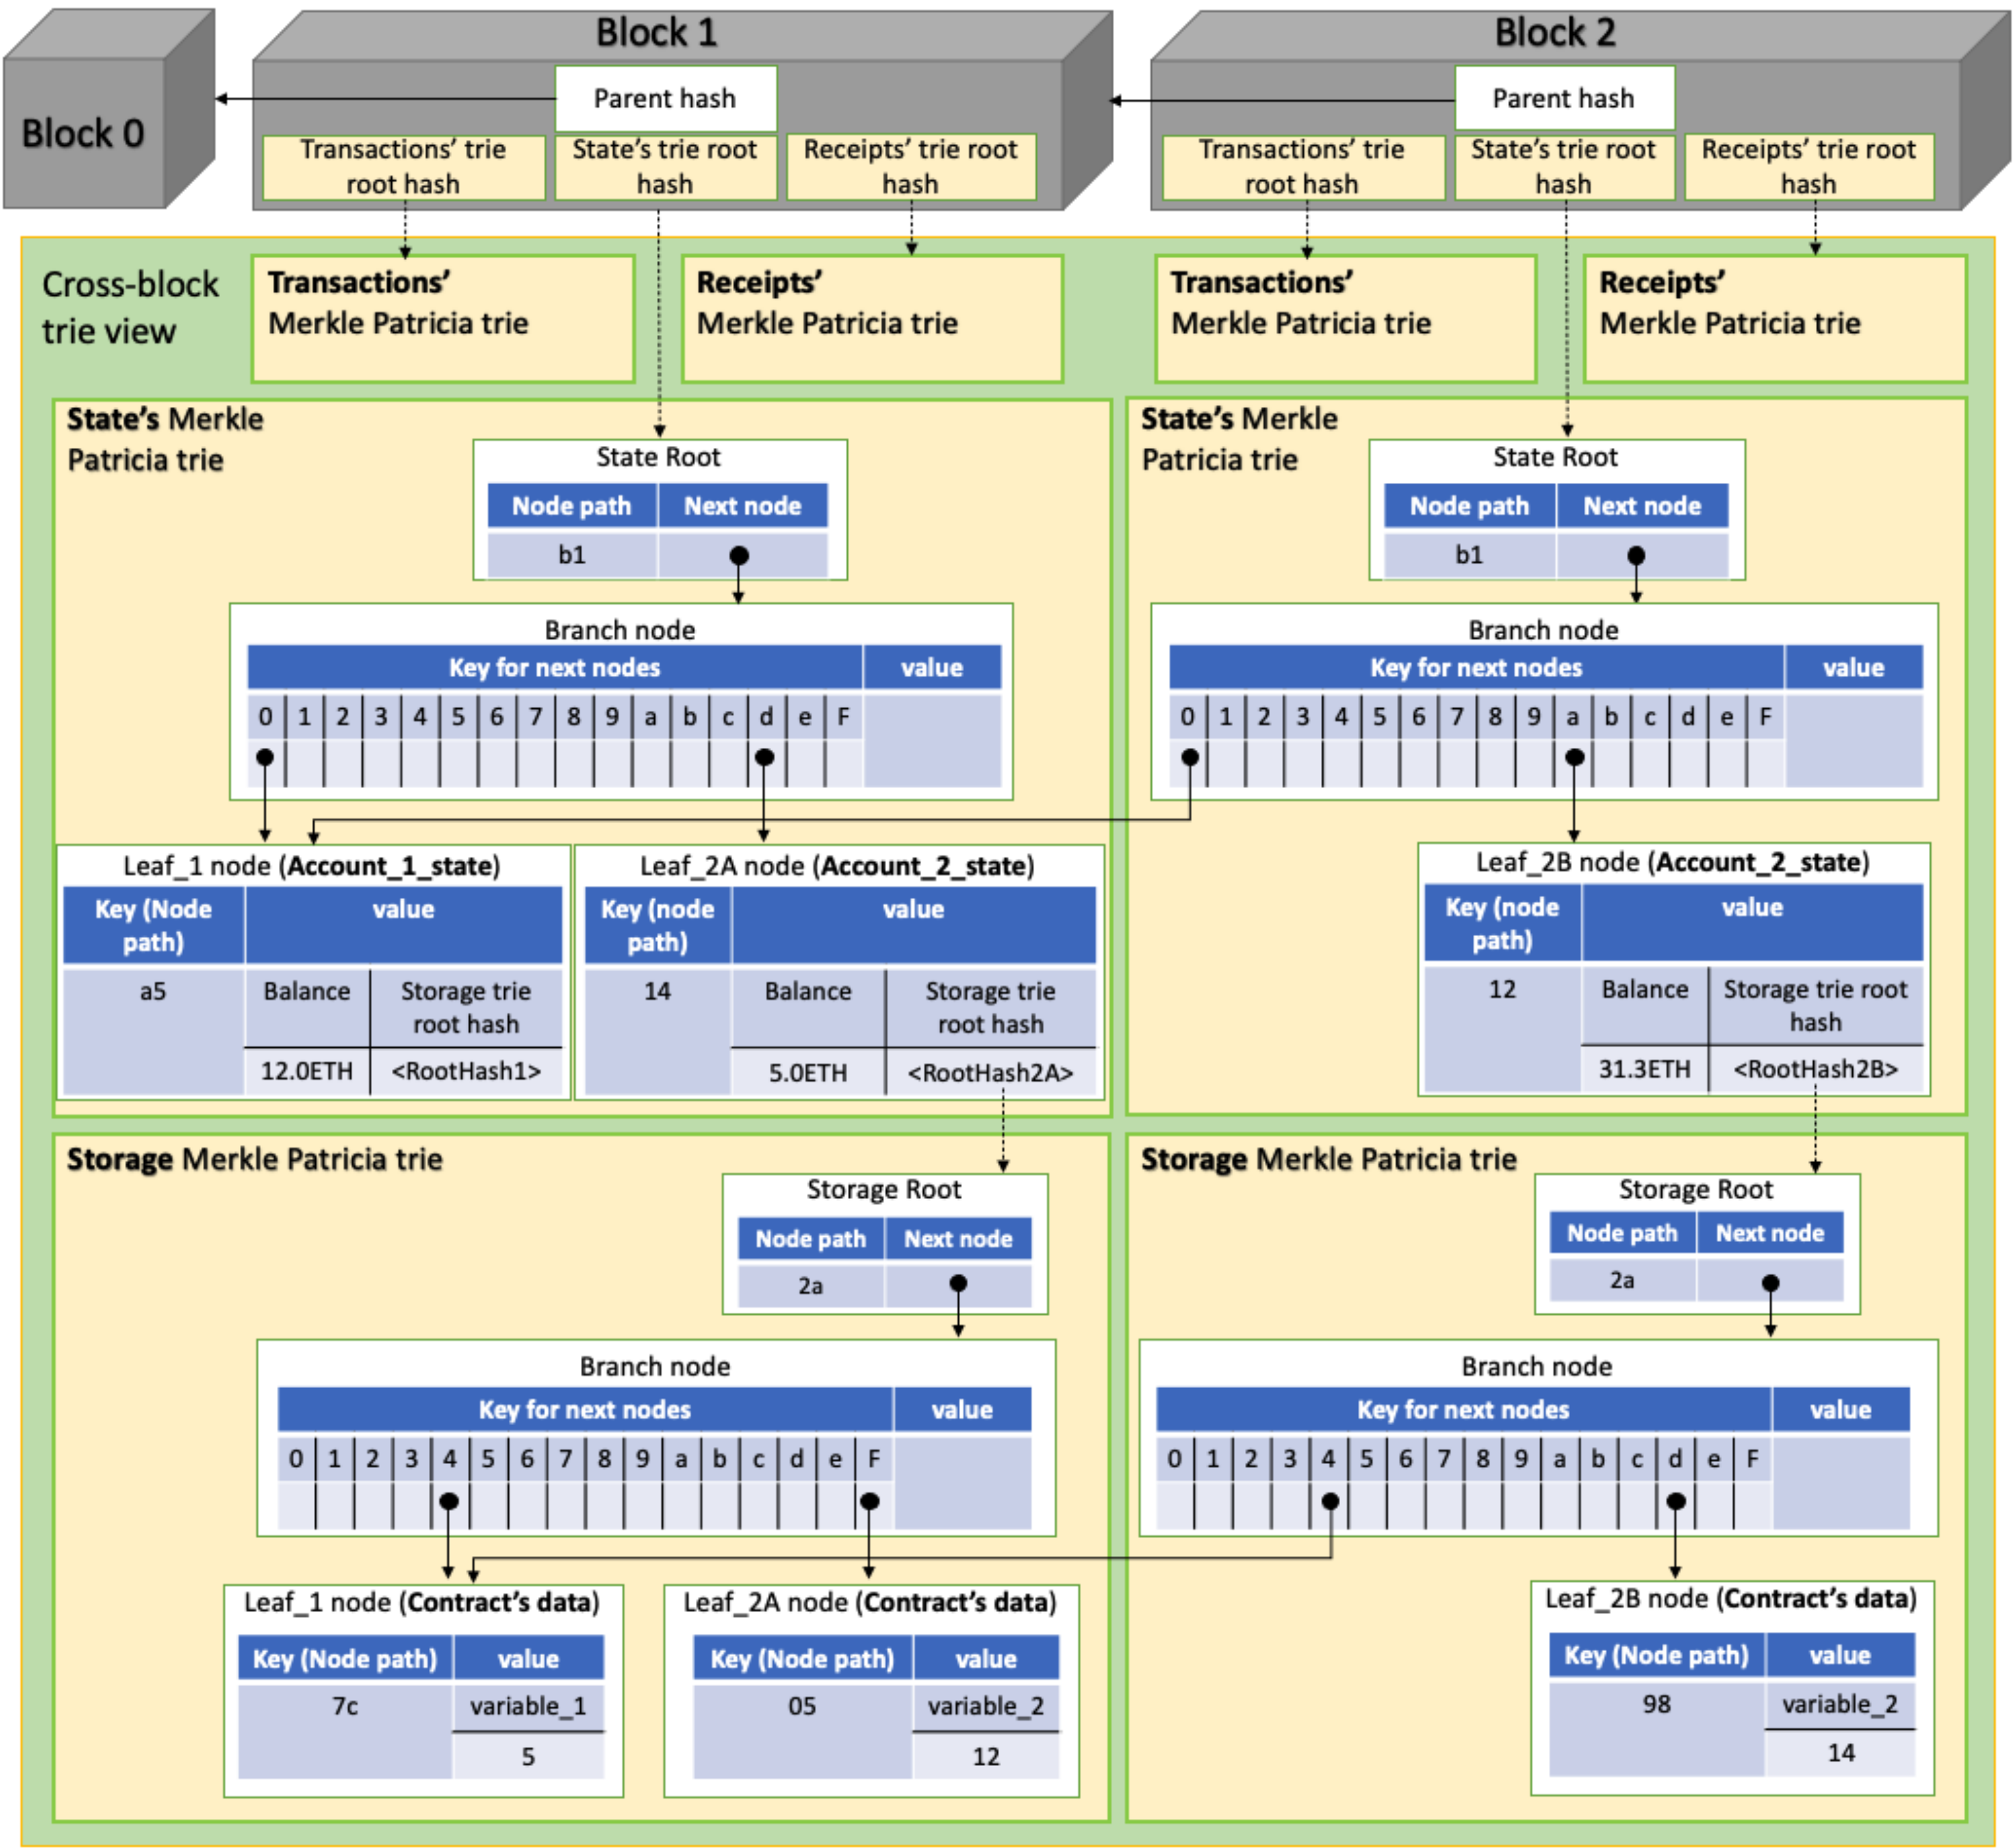
update  
lookup  
proof

StorageRoot is the root of a different trie.



# Ethereum

## State trie



# Ethereum

## Read contract state

1. ask trusted node
2. receive inclusion proof for

stateRoot: storageTrie  
account state: stateTrie

and block header



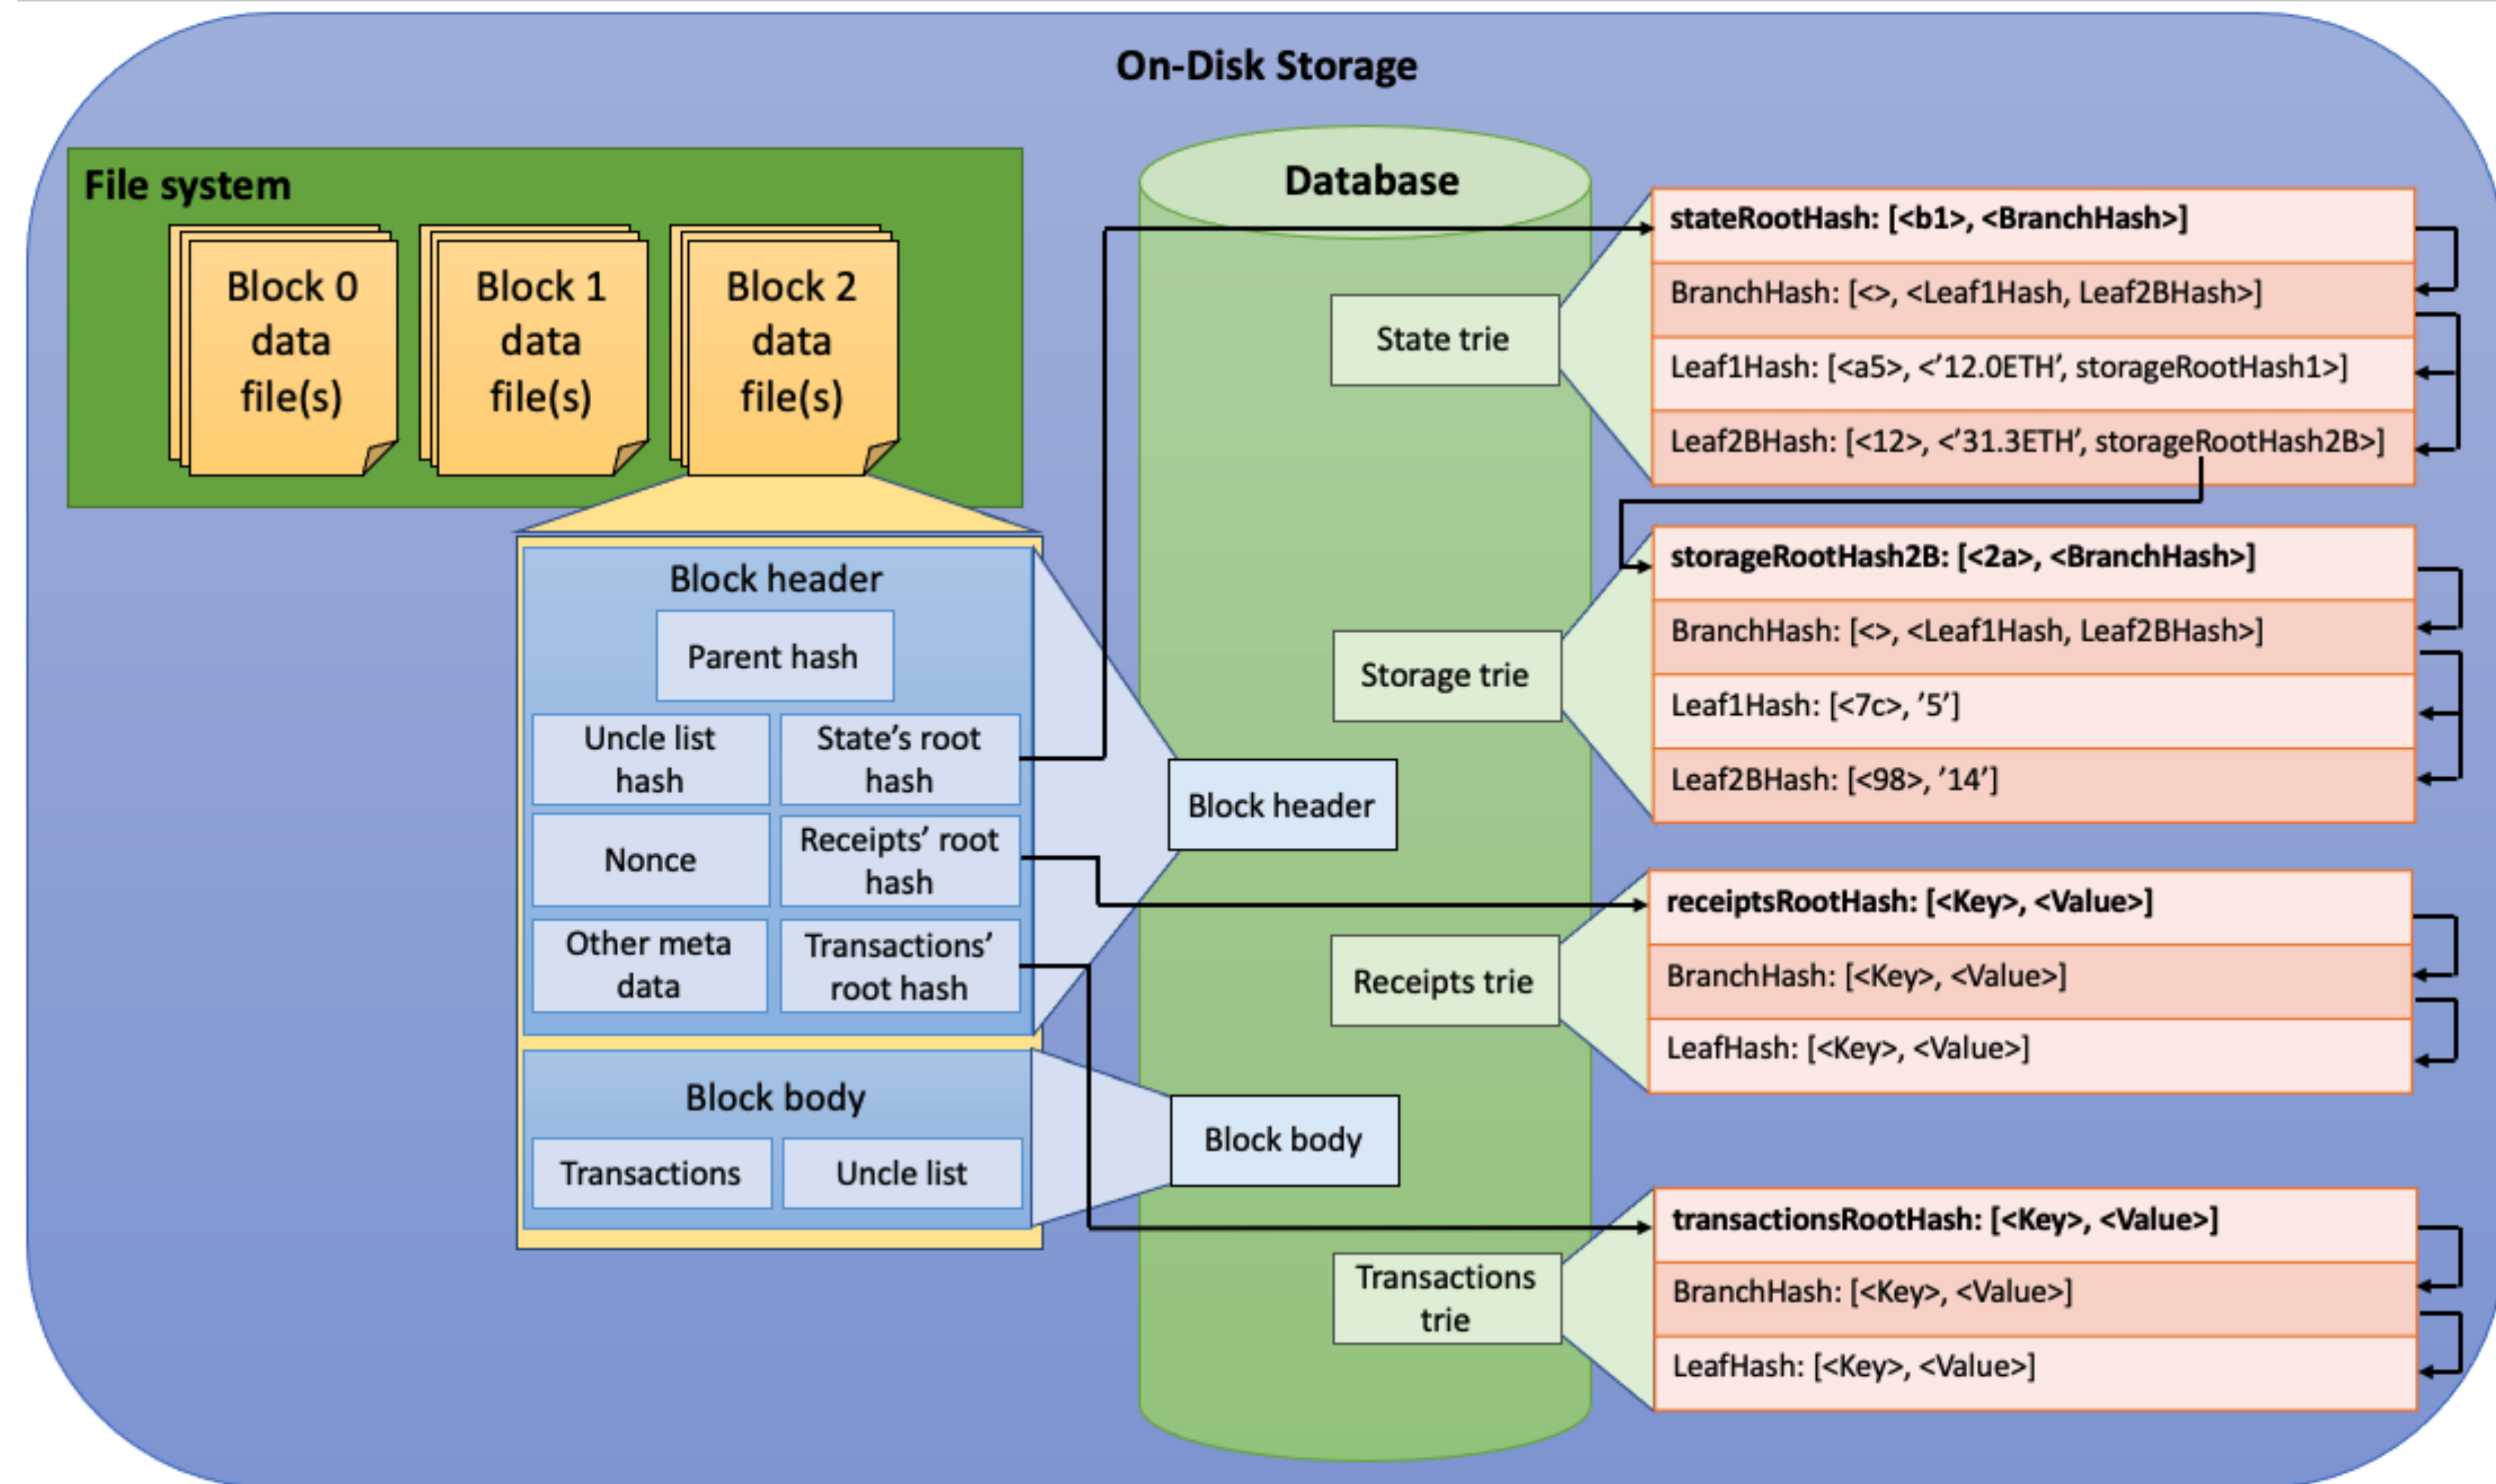
# Ethereum

## Receipts trie

Stores transaction results:

```
From: address
To: address
Status: ... // aborted?
Logs: events
ContractAddress address
// new contract address,
// if created
```

Return transaction results,  
by emitting Events,  
which are added to the logs.



# Ethereum

## Gas

How to pay transaction fees in Ethereum?

- all bytecode instructions have a cost specified in Gas
- transaction has fixed cost in Gas
- especially: storing values is expensive

Transactions specify *Gas price* and *Gas limit*

- *Gas price* is ether given per gas
- *Gas limit* is how much the transaction may spend at most

Actually: Gas price is divided by:  
Base price + Tip  
Base price is burned  
Tip is given to validators

# Ethereum

## Gas

Why specific gas per instruction:

- An infinite loop will cost infinitely much gas -> avoid denial of service

What happens if you hit the *Gas limit*?

- Exception is thrown and transaction reverted.
- Gas is still payed!

Which transactions are included?

- Miners will include transactions offering the highest gas price.
- Blocks have maximum amount of gas.

# Ethereum

## Gas - London upgrade 2021

Gas price is divided in Base price + Tip

- Base price is burned
- Tip is given to the validators

Blocks can be bigger than target size, but included transactions have to pay a larger base price